

# Low level documentation for API

## Overview

This documentation provides detailed technical specifications for the FastAPI-based backend service that powers the Dynamic Forecasting Visualization application. The service handles model training, prediction, drift detection, and monitoring.

## System Architecture

### Core Components

- 1. **FastAPI Application:** Main web service framework
- 2. **Forecasting Model:** Prophet-based time series forecasting
- 3. **Drift Detector:** Statistical concept drift detection
- 4. **MLflow Integration:** Model tracking and experimentation
- 5. **Prometheus Metrics:** System and application monitoring

## Configuration

### Key Configuration Parameters

Parameter	Description	Default Value
<code>BASE_DIR</code>	Application root directory	Path to parent directory
<code>MODEL_PATH</code>	Path to serialized model	<code>./models/model.pkl</code>
<code>METADATA_PATH</code>	Path to model metadata	<code>./models/metadata.json</code>
<code>SEASONAL_PERIODS</code>	Seasonality parameter	30
<code>PROMETHEUS_PORT</code>	Metrics endpoint port	8001
<code>MLFLOW_TRACKING_URI</code>	MLflow tracking server	<code>http://localhost:5001</code>
<code>MLFLOW_EXPERIMENT_NAME</code>	MLflow experiment name	<code>DemandForecasting</code>

## State Management

## Global Variables

Variable	Type	Description
<code>model</code>	<code>Optional[ForecastingModel]</code>	Loaded forecasting model
<code>df_global</code>	<code>Optional[pd.DataFrame]</code>	Loaded dataset
<code>drift_detector</code>	<code>Optional[DriftDetector]</code>	Drift detection instance

## Metrics System

### Prometheus Metrics

#### Endpoint Metrics

Metric Name	Type	Description	Labels
<code>model_init_requests_total</code>	Counter	Total /model_init requests	method, status_code
<code>model_init_request_latency_seconds</code>	Histogram	/model_init latency	method
<code>last_trained_date_requests_total</code>	Counter	Total /last_trained_date requests	method, status_code
<code>last_trained_date_request_latency_seconds</code>	Histogram	/last_trained_date latency	method
<code>predict_for_date_requests_total</code>	Counter	Total /predict_for_date requests	method, status_code
<code>predict_for_date_request_latency_seconds</code>	Histogram	/predict_for_date latency	method
<code>api_requests_per_ip_total</code>	Counter	Requests per client IP	client_ip, endpoint, method

#### System Metrics

Metric Name	Type	Description
<code>system_cpu_usage_percent</code>	Gauge	CPU utilization

<code>system_memory_usage_percent</code>	Gauge	Memory utilization
<code>system_disk_usage_percent</code>	Gauge	Disk utilization
<code>system_network_bytes_sent</code>	Gauge	Network bytes sent
<code>system_network_bytes_recv</code>	Gauge	Network bytes received
<code>system_file_handles</code>	Gauge	Open file handles

## Middleware

### `EndpointMetricsMiddleware`

- **Purpose:** Tracks request metrics for all endpoints
- **Behavior:**
  1. Records start time
  2. Processes request
  3. Updates relevant metrics:
    - Increments appropriate counter
    - Records latency
    - Tracks requests per IP

## Core Functions

### `train_initial_model(df: pd.DataFrame)`

- **Purpose:** Trains the first model instance
- **Parameters:**
  - `df`: Initial training data (first 6 months)
- **Behavior:**
  1. Trains Prophet model
  2. Saves model to disk
  3. Saves metadata

### `retrain_model(start_date, df, drift_metadata=None)`

- **Purpose:** Retrains model with updated data
- **Parameters:**
  - `start_date` : Starting date for training data
  - `df` : Complete dataset
  - `drift_metadata` : Optional drift detection info
- **Behavior:**
  1. Subsets data up to `start_date`
  2. Starts MLflow run with appropriate tags
  3. Trains new model
  4. Logs drift parameters if applicable

### `get_last_trained_date()`

- **Purpose:** Retrieves last training date from metadata
- **Returns:** `datetime` or `None`
- **Behavior:**
  1. Checks for metadata file
  2. Parses last trained date if exists

## API Endpoints

### `POST /model_init`

- **Purpose:** Initialize forecasting system
- **Parameters:**
  - `file` : CSV data file
  - `window_size` : Drift detection window
  - `threshold` : Drift detection threshold
- **Flow:**
  1. Loads and validates data

2. Initializes drift detector
3. Trains initial model
4. Returns success/error response

#### GET /last\_trained\_date

- **Purpose:** Get last model training date
- **Response:**

```
{  
  "last_trained_date": "YYYY-MM-DD"  
}
```

- **Error Handling:** Returns 404 if no model exists

#### GET /predict\_for\_date

- **Purpose:** Get prediction for specific date
- **Parameters:**
  - `current_date`: Date to predict (YYYY-MM-DD)
- **Response:**

```
{  
  "date": "YYYY-MM-DD",  
  "predicted": float,  
  "actual": float,  
  "error": float,  
  "drift_detected": bool  
}
```

- **Flow:**
  1. Validates date
  2. Generates prediction
  3. Checks for drift

4. Triggers retraining if drift detected

## Startup Sequence

1. Configures MLflow tracking
2. Starts Prometheus metrics server
3. Launches system metrics collector thread
4. Loads existing model if available

## Monitoring

- **Prometheus:** Scrapes metrics from `:8001`
- **MLflow:** Tracks all model training runs
- **System Metrics:** Collected every 5 seconds

## Error Handling

- All endpoints return 500 status for unexpected errors
- Detailed error messages logged
- MLflow tracks failed runs

## Dependencies

- FastAPI (web framework)
- Prophet (forecasting)
- Pandas (data processing)
- MLflow (experiment tracking)
- Prometheus Client (metrics)
- Psutil (system monitoring)